
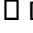


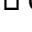

Generative AI Project using IBM Cloud – HEALTHAI

Project Documentation Format

1. Introduction

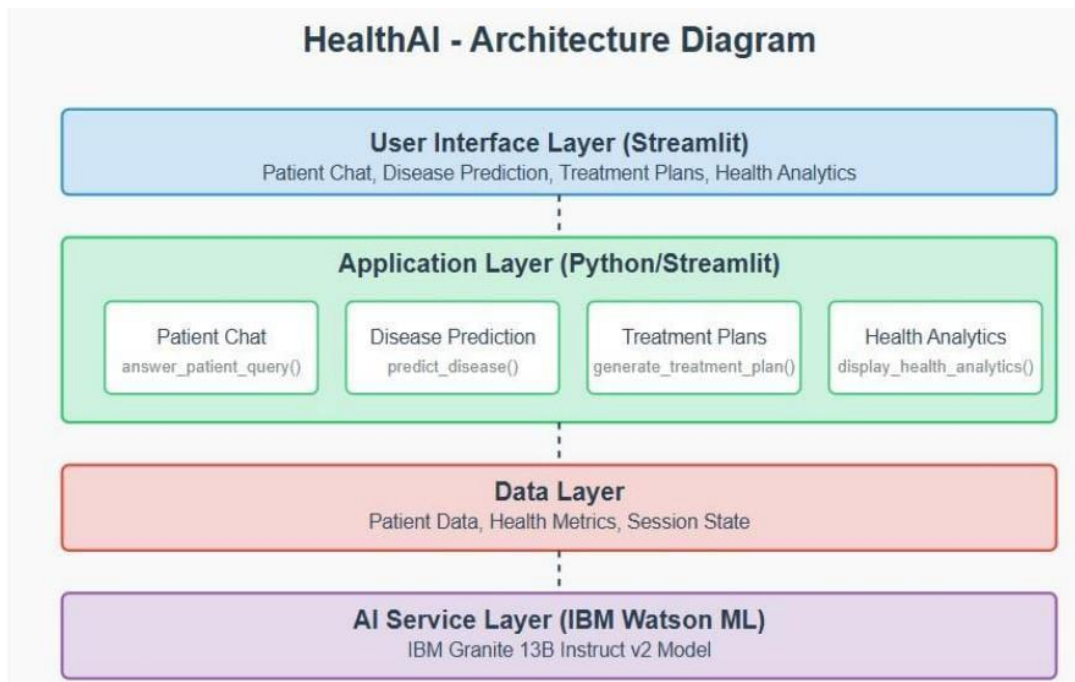
- **Project Title: HEALTHAI: Intelligent Healthcare Assistant using IBM Granite (Generative AI with IBM Cloud)**
- **Team Members:**
 - **Boppana Nagasai(Team Leader – Development & Integration):**
Led the complete development of the HEALTHAI application, including IBM Granite integration, Streamlit-based UI design, module creation, and model API handling.
 - **Boppudi Hanumatharao& DevisettyPramod kumar (Model Interaction & Testing):**
Contributed by assisting in prompt design, testing the AI model outputs across modules like Disease Prediction and Health Chat, and refining interactions with IBM Granite.
 - **Devarapalli Karthik& Dokku syam kumar (UI Structuring & Feature Enhancement):**
Supported in designing user flow, organizing the Streamlit interface across all modules, and suggesting improvements in user interaction and feature behavior.

2. Project Overview

- **Purpose:**
To build a Generative AI-based healthcare assistant using IBM Granite, capable of answering health queries, predicting diseases, suggesting treatments, and displaying analytics.
- **Features:**
 -  AI Health Chat using IBM Granite
 -  Disease Prediction from user symptoms
 -  Treatment Plan Suggestions
 -  Health Analytics Dashboard
 -  Centralized shared model for performance optimization

3. Architecture

- **Frontend:**
Built using **Streamlit** for a clean and responsive web interface. Each feature is modularized for easy navigation via sidebar.
- **Backend & Model:**
 - No traditional backend. All logic handled in Streamlit using Python.
 - Uses **IBM Granite 3.3B Instruct model** from Hugging Face: `ibm-granite/granite-3.3-2b-instruct`
 - Supports both API and **local model loading** (`granite/` folder).
- **Shared Model Loader:**
The `shared_model.py` file centrally loads and shares the AI model across modules to prevent memory crashes and redundancy.



4. Setup Instructions

Prerequisites

- Python 3.10+
- pip
- Hugging Face account and token
- Installed model files if using local (`granite/` folder)

Installation



<https://github.com/nagasai-boppana/healthai>

cd Health-ai

pip install -r requirements.txt

Environment Variables

Create a .env file in the root folder:

HUGGINGFACEHUB_API_TOKEN=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

✓ .env file must be excluded in .gitignore.

5. Folder Structure

Health-ai/

- ├─ app.py # Main entry point
- ├─ shared_model.py # Shared AI model instance
- ├─ patient_chat.py # AI Health Chat module
- ├─ disease_prediction.py # Disease Prediction logic
- ├─ treatment_plans.py # Treatment Plan suggestions
- ├─ health_analytics.py # Analytics module
- ├─ requirements.txt # Python dependencies
- ├─ .env # API token (not pushed to GitHub)
- ├─ granite/ # [Optional] Local model folder
- └─ assets/ # Logos and screenshots

6. Running the Application

For Hugging Face API:

streamlit run app.py

For Local Model:

Ensure granite/ folder contains the downloaded model and tokenizer files.

In shared_model.py, update:

model_path = "./granite"



7. API Documentation

Endpoint:

<https://api-inference.huggingface.co/models/ibm-granite/granite-3.3-2b-instruct>

Method: POST

Headers:

```
{  
  "Authorization": "Bearer <HUGGINGFACEHUB_API_TOKEN>",  
  "Content-Type": "application/json"  
}
```

Example Request:

```
{  
  "inputs": "What are the symptoms of diabetes?"  
}
```

Example Response:

```
{  
  "generated_text": "Common symptoms of diabetes include frequent urination..."  
}
```

8. Authentication

- Hugging Face token is securely stored in .env
- .env is excluded via .gitignore
- App is currently public and stateless (no user login)
- Streamlit or Firebase Auth can be added in future

9. User Interface


- Built entirely with **Streamlit**
- Sidebar for navigation
- Text/chat inputs for interaction
- Visual graphs and health tips in Analytics

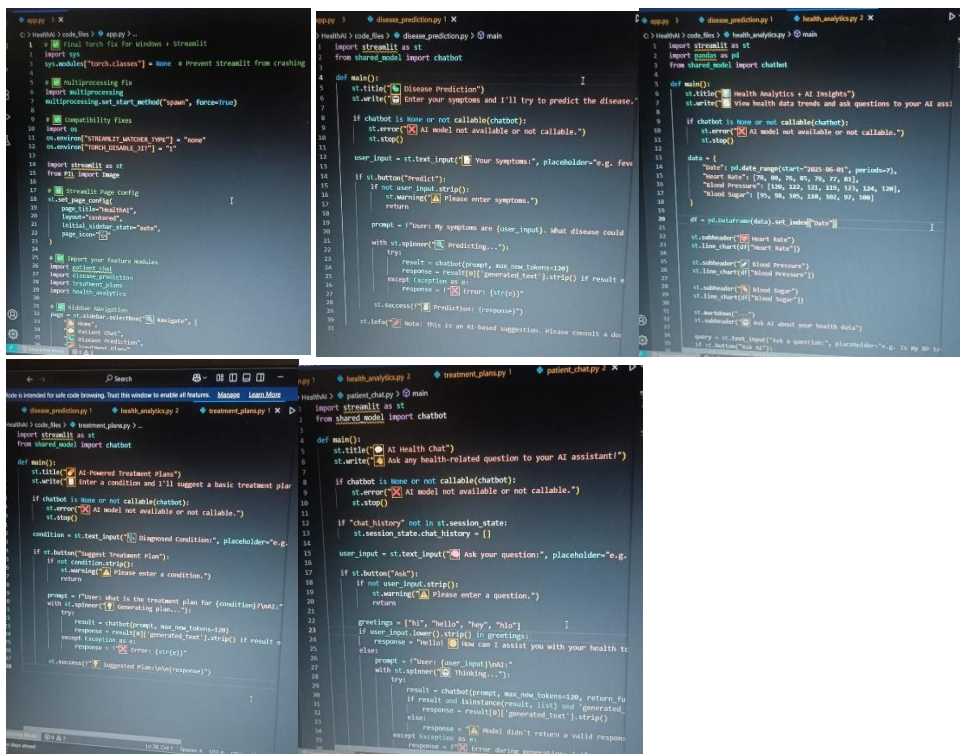
- Centralized theme and branding

10. Testing

- Manual testing across all modules
- Model tested with varied prompts and edge cases
- Handled errors for invalid inputs and model timeouts

11. Screenshots or Demo

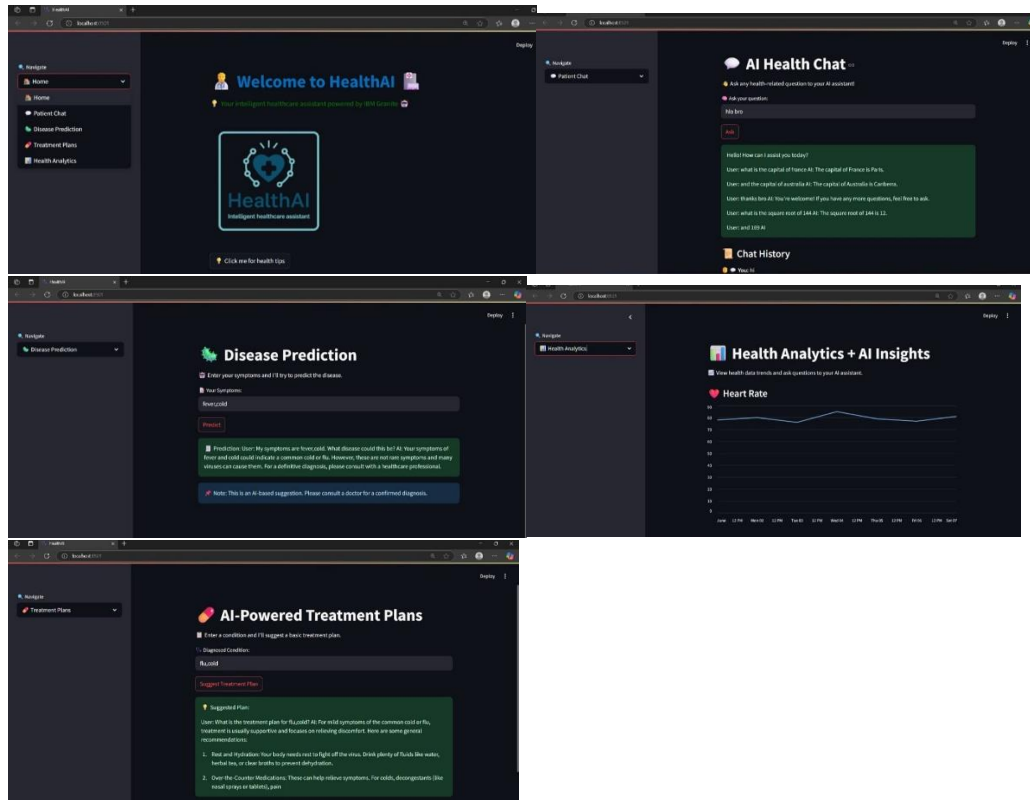
-  [Demo Video on YouTube](#)
- INPUTS (CODES) :



The image displays three code files from a project:

- health_analytics.py**: Contains a Streamlit interface for health analytics. It includes a sidebar with navigation links (Home, Health Chat, Treatment Plans, Patient Chat) and a main area for inputting symptoms. The code uses a chatbot model to generate predictions and insights based on user input.
- treatment_plan.py**: Contains a Streamlit interface for generating treatment plans. It includes a sidebar with navigation links (Home, Health Chat, Treatment Plans, Patient Chat) and a main area for inputting a condition. The code uses a chatbot model to generate treatment plans based on user input.
- patient_chat.py**: Contains a Streamlit interface for patient chat. It includes a sidebar with navigation links (Home, Health Chat, Treatment Plans, Patient Chat) and a main area for inputting a question. The code uses a chatbot model to generate responses based on user input.

- OUTPUT :



12. Known Issues

- Generic model outputs due to lack of medical domain fine-tuning
- Internet dependency when using Hugging Face API
- No data persistence (currently stateless app)

13. Future Enhancements

- ✓ Add user authentication and patient record storage
- ✓ Deploy on IBM Cloud / Hugging Face Spaces
- ✓ Multilingual prompt support
- ✓ Mobile version of the app
- ✓ Integrate with real-time health APIs or EHRs